

# DISSANAYAKE E.G.M — Deployment & Evaluation Engineer

This version will be:

- ✓ Fully aligned with the *actual state* of SmartOps (your orchestrator, telemetry stack, ERP simulator, chaos features, RBAC status).
  - ✓ Structured professionally like an industry SRE / DevOps implementation plan.
  - ✓ No illusions — only tasks that match what exists or is achievable within your architecture.
  - ✓ Clean, clear, and ready to paste into your project report.
- 

# DISSANAYAKE E.G.M — COMPREHENSIVE IMPLEMENTATION PLAN

## **Role: Deployment, Observability & Resilience Validation Lead**

Dissanayake owns the **infrastructure backbone** of SmartOps — the Kubernetes platform, CI/CD automation, telemetry stack, chaos testing suite, and evaluation pipelines. His responsibility is to ensure SmartOps can be **deployed, observed, stress-tested, and benchmarked reliably** in a real cluster.

Workload share: **27–30%**

---

## **SECTION 1 — ROLE OVERVIEW & OBJECTIVE**

Dissanayake is responsible for:

- ✓ **Building and maintaining the Kubernetes platform**
- ✓ **Implementing CI/CD workflows**
- ✓ **Deploying all SmartOps services (orchestrator, ERP simulator, AI agents, policy engine)**
- ✓ **Integrating telemetry: Prometheus, Loki, Grafana, OTel Collector, Tempo**

✓ **Installing and running chaos experiments**

✓ **Producing MTTR, SLO, SLA, recovery KPIs**

✓ **Final benchmarking for viva presentation**

He ensures the system operates *continuously*, auto-heals under stress, and provides complete visibility from metrics to logs to traces.

---

## SECTION 2 — ENVIRONMENT & PLATFORM SETUP (WEEKS 1–2)

### Tasks

- Create GitHub organization (smartops-ai) and configure:
  - CODEOWNERS
  - main branch protection
  - PR templates
  - issue templates
  - Actions secrets (KUBECONFIG\_DEV, GHCR\_PAT, etc.)
- Provision Kubernetes cluster (Docker Desktop / Minikube / K3s / cloud)
- Create namespaces:

```
smartops-dev  
smartops-stage  
smartops-prod
```

- Initialize the Helm umbrella chart in:

```
platform/helm/smartops/
```

containing:

- Prometheus stack
- Loki
- Tempo
- OTel Collector
- ERP Simulator chart
- Orchestrator chart (already exists)

### Collaboration

- Peiris → align orchestrator RBAC & service account permissions
- Team → ensure secrets and endpoints shared

## Acceptance

- ✓ Helm install succeeds
  - ✓ Grafana dashboard accessible
  - ✓ CI skeleton workflow runs
- 

# SECTION 3 — CI/CD PIPELINE IMPLEMENTATION (WEEKS 1–3)

## CI/CD Responsibilities

- Build Docker images → Push to GHCR
- Deploy orchestrator, ERP simulator, AI agents, policy-engine
- Automatic promotion workflow (dev → stage → prod)
- Rollback capability

## Tasks

- Implement reusable GitHub Actions:

```
ci.yml  
deploy-dev.yml  
promote-stage.yml  
promote-prod.yml
```

- Add artifact caching, parallel jobs
- Validate Helm lint + Kubernetes manifest tests
- Configure semantic versioning (v0.1.0, v0.1.1, etc.)
- Add rollback logic:

```
helm upgrade --install smartops . --atomic --wait
```

## Collaboration

- Kulathunga → integrate model unit tests, model registry build
- Gunarathne → policy engine test suite
- Peiris → orchestrator build + startup validation

## Acceptance

- ✓ Merge to `main` → auto deploys DEV
  - ✓ Tag release → deploys to STAGE
  - ✓ Rollback works via GitHub Actions
-

# **SECTION 4 — TELEMETRY & OBSERVABILITY STACK (WEEKS 2–4)**

## **Deploy Observability Components**

- Prometheus Operator
- Loki Stack (log aggregation)
- OTel Collector (trace pipeline)
- Tempo (distributed tracing backend)
- Grafana (visualization layer)

## **Configure Data Sources**

- Prometheus scrape targets:
  - orchestrator
  - ERP simulator
  - future AI services
  - kube-state-metrics
- Loki logs:
  - orchestrator logs
  - pod crash logs
- Tempo traces via OTLP
- Dashboards using JSON or GUI

## **Collaboration**

- Peiris → validate orchestrator metrics & tracing
- Kulathunga → export AI agent metrics
- Gunarathne → add policy engine metrics

## **Acceptance**

- ✓ “MTTR”, “Error rate”, “Closed-Loop Actions” dashboards working
  - ✓ Logs + metrics + traces linkable by trace ID
  - ✓ Alerts firing: pod restart, SLA breach, anomaly spike
- 

# **SECTION 5 — CHAOS ENGINEERING SETUP (WEEKS 3–4)**

## **Chaos Tools**

- Chaos Mesh
  - (Or Chaos Toolkit depending on resource availability)

## Experiments Owned

- Podkill: kill orchestrator or ERP simulator pod
- CPU stress: 80–100% CPU
- Memory leak: 500–800MB leak
- Network delay: 200–1000ms
- Packet loss: 10–30%
- Combined scenarios: multiple chaos conditions

## How chaos supports SmartOps

Chaos validates:

- AI anomaly detection
- RCA mapping
- Closed-loop actions
- Kubernetes recovery
- MTTR metrics

## Collaboration

- Peiris → orchestrator failure scopes
- Kulathunga → training data windows for chaos
- Gunarathne → policy rule interaction under chaos

## Acceptance

- ✓ All chaos experiments execute safely
  - ✓ Recovery is measurable (MTTR shown on Grafana)
  - ✓ Logs stored for evaluation
- 

# SECTION 6 — DEPLOYMENT & RELEASE MANAGEMENT (WEEKS 4–10)

## Responsibilities

- Maintain Helm values per environment
- Ensure safe promotions:
  - dev → stage → prod
- Validate Kubernetes health gates
- Maintain image version chart in GitHub Actions

## Promotion Policy

- Stage deploy only allowed if:
  - No failing workloads
  - Prometheus alerts in green
  - MTTR < threshold
- Prod deploy only allowed if:
  - Stage “golden window” of 72 hours passes
  - No anomaly spikes
  - No persistent alerts

## Collaboration

- Gunarathne → deploy policy-engine subchart
- Peiris → ensure orchestrator scales/restarts correctly

## Acceptance

- ✓ Prod deployment successful with no failing pods
  - ✓ Rollout verified with Grafana + Prometheus
- 

# SECTION 7 — VISUALIZATION & OPS CONSOLE (WEEKS 8–14)

## Dashboards Required

1. **Cluster Overview**
  - CPU, RAM, nodes, pod restarts
2. **Orchestrator Actions Panel**
  - scale, restart, patch
  - closed-loop duration
  - queue depth
3. **AI Agent Metrics**
  - anomaly scores
  - RCA confidence
  - model drift
4. **Policy Engine Metrics**
  - guardrail hits
  - policy evaluation latency
  - approved vs blocked actions
5. **Chaos vs Auto-Heal Timeline**
  - chaos event
  - anomaly detection
  - RCA
  - orchestrator action
  - verification

## Alerts

- MTTR > threshold
- CrashLoopBackOff
- Missing metrics from orchestrator
- AI inference latency too high
- Policy conflicts too frequent

## Acceptance

- ✓ Dashboards auto-refresh
  - ✓ Alerts trigger within <60s
  - ✓ All components are observable
- 

# SECTION 8 — RESILIENCE EVALUATION & BENCHMARKING (WEEKS 11–14)

## Design Evaluation Test Matrix

Dimensions:

- Chaos Type
- AI Model Variant
- Policy Engine Rule Set
- Workload Intensity

## Metrics to Capture

- MTTR reduction percentage
- Auto-heal success rate
- Correct vs incorrect RCA
- SLO adherence
- SLA uptime
- Closed-loop duration
- Action success vs failure
- Number of guardrail blocks

## Output

Produce:

/tests/chaos/results/  
/docs/reports/evaluation.pdf

## **Acceptance**

- ✓ MTTR improved  $\geq 40\%$
  - ✓ Auto-heal  $\geq 90\%$  accuracy
  - ✓ SLA  $\geq 99\%$  maintained
- 

# **SECTION 9 — HARDENING & SECURITY (WEEKS 14–16)**

## **Tasks**

- Tighten RBAC for orchestrator, AI agents, policy engine
- Apply network policies limiting:
  - namespace egress
  - cross-service calls
- Add resource limits & HPA targets
- Enable image signing (cosign)
- Add vulnerability scanning
- Add promotion SLO gates (cannot deploy if SLO  $< 98\%$ )

## **Acceptance**

- ✓ No unauthorized API paths above allowed verbs
  - ✓ 72-hour staging observation with no critical alerts
  - ✓ Prod promotions gated by SLO score
- 

# **SECTION 10 — DOCUMENTATION & KNOWLEDGE TRANSFER**

## **Must Produce**

- /docs/runbooks/\*
- /docs/deployment\_guide.md
- /docs/chaos\_manual.md
- “How to Reproduce Evaluation” guide
- Grafana dashboard index

## **Acceptance**

- ✓ All docs reproducible
  - ✓ Fresh-cluster deploy validated
  - ✓ Viva-ready
- 

## SECTION 11 — KPIs (FOR DISSANAYAKE'S EVALUATION)

KPI	Target	Tool
Deployment success rate	$\geq 95\%$	GitHub Actions
Chaos test reliability	$\geq 90\%$	Chaos Mesh
MTTR reduction	$\geq 40\%$	Grafana MTTR panel
Dashboard accuracy	$\geq 90\%$	Grafana vs Loki logs
SLA adherence	$\geq 99\%$	Prometheus Query

---

## SECTION 12 — INTERFACES & TEAM COLLABORATION

Collaborator	Interaction
<b>Peiris P.V.G</b>	consumes logs/metrics; validates orchestrator actions under chaos
<b>Kulathunga K.A.K.M</b>	supplies AI agent metrics and model performance
<b>Gunarathne M.D.C.H</b>	integrates policy metrics, guardrails, action plans
<b>Supervisor</b>	final evaluation & deployment quality assurance

---

## SECTION 13 — FINAL DELIVERABLES (OWNED ENTIRELY BY DISSANAYAKE)

1. **Helm Umbrella Chart & Values**
  2. **CI/CD Workflow Stack**
  3. **Chaos Testing Suite**
  4. **Grafana Dashboards**
  5. **Evaluation Reports**
  6. **Runbooks & Deployment Guides**
  7. **Stage/Prod Hardening Checklist**
-

## FINAL SUMMARY

Dissanayake ensures SmartOps is:

- Deployable
- Observable
- Chaos-validated
- Benchmarkable
- Secure
- Production-ready

His contributions form the **operational backbone** allowing Peiris (orchestrator), Kulathunga (AI), and Gunarathne (policy engine) to deliver a fully autonomous, observable, and resilient platform.