

# KULATHUNGA K.A.K.M — Lead AI Architect (Anomaly Detection & RCA)

Fully aligned with the *actual* current SmartOps architecture, the orchestrator capabilities Peiris implemented, the policy engine design, and the real telemetry stack already deployed.

---

## KULATHUNGA K.A.K.M — COMPREHENSIVE IMPLEMENTATION PLAN

(Clean, restructured, realistic, production-grade)

---

## SECTION 1 — ROLE OVERVIEW & OBJECTIVE

Kulathunga owns **all AI intelligence inside SmartOps**, specifically:

**Two core AI agents:**

1. **Anomaly Detection Agent**
  - Multivariate, real-time inference
  - Maps telemetry → anomaly type, score, severity
  - Sends structured signals into orchestrator
2. **RCA (Root Cause Analysis) Agent**
  - Inspects anomalies + dependency graphs
  - Produces top-ranked root causes with confidence
  - Produces human-readable explanations

**Why his role is critical:**

Every autonomous recovery in SmartOps starts with **AI signals** sent to:

/v1/signals/anomaly  
/v1/signals/rca

...which power the closed-loop orchestrator created by Peiris.

Kulathunga must ensure:

- ✓ AI signals are accurate
- ✓ RCA is explainable
- ✓ Real-time inference is stable
- ✓ Agents integrate smoothly with Orchestrator + Policy Engine
- ✓ Models learn from successful/failed actions

Workload Share: **25–27%**

Focus: **AI intelligence, ML pipelines, diagnostics, explainability, evaluation**

---

## SECTION 2 — SYSTEM CONTEXT (REAL PROJECT STATE)

SmartOps currently has:

- Peiris's orchestrator with:
  - scale/restart/patch
  - closed-loop manager
  - verification engine
  - complete Kubernetes API integration
- ERP simulator generating chaos modes
- Prometheus Operator, Loki, OTel Collector, Tempo
- Orchestrator already validating anomaly + RCA formats
- Policy Engine design underway

Kulathunga's AI agents plug in at the earliest stage of the automation loop:

Telemetry → AI Detect → AI RCA → Policy Engine → Orchestrator → K8s

---

## SECTION 3 — HIGH-LEVEL DELIVERABLES

Kulathunga must deliver:

1. **Anomaly Detection Agent**
2. **RCA Diagnostic Agent**
3. **Feature store & telemetry ingestion pipeline**
4. **Model training, tuning & evaluation notebooks**
5. **Chaos validation dataset**
6. **Prometheus + OTel metrics for AI agents**
7. **REST APIs for real-time inference**
8. **Feedback-driven retraining system**
9. **Full documentation bundle for AI side**

Everything below is structured to produce those deliverables reliably.

---

## SECTION 4 — PHASED IMPLEMENTATION PLAN (BEST VERSION)

---

### PHASE 1 — AI Repository & Environment Setup (Week 1–2)

#### Goals

Create a clean and scalable AI workspace.

#### Tasks

- Create folder structure:

```
apps/agent-detect/  
apps/agent-diagnose/  
  src/  
  models/  
  notebooks/  
  pipelines/
```

- Prepare runtime:
  - Python 3.10/3.11
  - PyTorch/TensorFlow (choose one early)
  - scikit-learn
  - Pandas
  - Uvicorn/FastAPI
- Create:
  - requirements.txt
  - Dockerfile
  - Makefile for train/test/infer
  - CI pipeline for lint + unit tests

#### Deliverables

- ✓ Repo builds successfully
- ✓ Unit test harness
- ✓ CI runs on PRs

---

## PHASE 2 — Telemetry Ingestion & Feature Engineering (Week 2–4)

### Telemetry Sources

- Prometheus metrics
- Loki logs (JSON)
- Tempo traces (optional later)
- ERP Simulator chaos modes

### Feature Store Requirements

- Extract fixed-length multivariate windows (30s, 60s, 120s)
- Features:
  - CPU usage
  - Memory usage
  - Error rate
  - Latency and p95/p99
  - Jitter
  - Request throughput
  - Chaos mode indicators
- Store as parquet in:

/data/processed/

### Tasks

- Build Prometheus query pipeline (Python HTTP client)
- Build log sampler from Loki
- Normalize time windows
- Label “normal” vs “faulted” windows using ERP chaos

### Deliverables

- ✓ Dataset ready
- ✓ Feature schema documented
- ✓ Telemetry aligned with orchestrator signal format

---

## PHASE 3 — Anomaly Detection Model (Week 4–7)

### Model Requirements

- Real-time inference (<3–5 seconds)

- Accurate:  $F1 \geq 0.85$
- Low false positives

## Architecture

- LSTM Autoencoder
- Optional: GRU encoder for low latency
- Adaptive thresholding:
  - Rolling mean
  - MAD
  - z-score
  - Combined threshold logic

## Outputs

Agent must return JSON:

```
{
  "windowId": "...",
  "service": "erp-simulator",
  "isAnomaly": true/false,
  "score": 0.87,
  "type": "latency" | "resource" | "error" | "cpu" | ...
}
```

## Deliverables

- ✓ Model saved under /models/
  - ✓ Inference engine with ~5s per request
  - ✓ Chaos dataset validation
- 

# PHASE 4 — RCA (Root Cause Analysis) Engine (Week 6–9)

## Approach (Realistic + Lightweight)

- Use dependency graph:
  - ERP → DB → Cache → Worker → API
- Compute deltas in metrics:
  - CPU spike
  - Memory leak
  - Error burst
  - Latency jitter
- Assign probabilities
- Provide explainable output

## Output Format

```
{  
  "windowId": "...",  
  "service": "erp-simulator",  
  "rankedCauses": [  
    {"svc": "erp-simulator", "cause": "memory_leak", "probability": 0.92}  
  ],  
  "confidence": 0.88  
}
```

## Deliverables

- ✓ RCA accuracy  $\geq 80\%$
  - ✓ Explainability module (rules + metrics-based)
  - ✓ Works on dev cluster chaos tests
- 

# PHASE 5 — Model Evaluation & Benchmarking (Week 8–10)

## Chaos-based evaluation

Use ERP Simulator:

- /chaos/memory-leak/enable
- /chaos/cpu-spike/enable
- /chaos/latency-jitter/enable
- /chaos/error-burst/enable

## KPIs

- F1 score
- False positive rate
- Latency per inference
- RCA confidence

## Deliverables

- ✓ /docs/reports/model\_eval.pdf
  - ✓ Confusion matrices
  - ✓ Comparison between classical vs deep models
- 

# PHASE 6 — API Integration & Service Deployment (Week 9–11)

## Expose APIs

Two microservices:

### **1. agent-detect**

POST /v1/anomaly/predict

### **2. agent-diagnose**

POST /v1/rca/diagnose

## **Requirements**

- Respond within <5 seconds
- Add Prometheus metrics:
  - anomaly\_inference\_latency\_seconds
  - anomaly\_requests\_total
  - rca\_inference\_latency\_seconds

## **Integration**

- Send anomaly → /v1/signals/anomaly
- Send RCA → /v1/signals/rca

## **Deliverables**

- ✓ Both AI services callable inside cluster
  - ✓ Orchestrator integration validated
- 

# **PHASE 7 — Continuous Learning Loop (Week 11–13)**

## **Data Sources**

Use orchestrator outputs:

- Verification success
- Action type
- Closed-loop duration
- Deployment status

## **Tasks**

- Build feedback DB
- Tag windows with “success/failure”
- Retrain anomaly & RCA models periodically
- Track drift

## **Deliverables**

- ✓ Retraining workflow
  - ✓ Drift detection dashboard
  - ✓ Improved MTTR after learning
- 

## PHASE 8 — Security, Governance & Deployment (Week 13–15)

### RBAC

AI agents should only be allowed to:

- Read telemetry
- NOT directly call Kubernetes
- NOT modify cluster resources

All orchestration must go through Peiris's orchestrator.

### Deployment Requirements

- Helm subchart
- Resource limits
- Readiness & liveness probes
- Model artifacts stored safely

### Deliverables

- ✓ RBAC YAML
  - ✓ Helm chart for AI services
  - ✓ GitHub Actions deploy pipeline
- 

## PHASE 9 — Documentation & Viva Preparedness (Week 15–16)

### Documents to Produce

- ai\_agent\_architecture.md
- model\_training\_guide.md
- model\_eval.pdf
- api\_reference.md
- telemetry\_feature\_engineering.md
- rca\_design.md

## Acceptance

- ✓ All AI docs merged to /docs/ai/
  - ✓ Review-ready for examiner
- 

## SECTION 5 — KPIs (REALISTIC TARGETS)

Metric	Target	Validation
F1 Score	$\geq 0.85$	Evaluation notebooks
False Positive Rate	$\leq 2\%$	Chaos dataset
RCA Confidence	$\geq 80\%$	RCA logs
Inference Latency	$\leq 5$ sec	Prometheus metrics
Auto-Heal MTTR Improvement	$\geq 30\text{--}40\%$	Grafana panels
Model Drift Detection	< 24h	Drift monitor

---

## SECTION 6 — COLLABORATION INTERFACES

### With Peiris (Orchestrator)

- Validate anomaly & RCA JSON schemas
- Provide signals that map directly to closed-loop actions
- Ensure inference latency matches closed-loop rate

### With Gunarathne (Policy Engine)

- Provide anomaly types and RCA causes mapped to DSL keywords
- Provide confidence thresholds
- Respond to policy-engine feedback loop

### With Dissanayake (Infra/DevOps)

- Deploy AI agents
  - Configure Grafana dashboards
  - Enable data scraping from Prometheus
-

# SECTION 7 — FINAL DELIVERABLES

1. **Anomaly Detection Service** (/apps/agent-detect/)
  2. **RCA Diagnostic Service** (/apps/agent-diagnose/)
  3. **AI Models & Registry** (/models/)
  4. **Feature Store & ETL Pipelines** (/data/processed/)
  5. **Evaluation Reports** (/docs/reports/)
  6. **AI Documentation Set** (/docs/ai/)
  7. **Helm Charts & CI Deployments** (platform/helm/smartops/agent-\*)
- 

## SUMMARY — Kulathunga's Role in SmartOps

Kulathunga builds the **intelligence** that allows SmartOps to:

- Detect failures
- Diagnose root causes
- Trigger autonomous remediation
- Learn from actions
- Improve MTTR continuously

He is:

- **Accountable** for all ML components
- **Responsible** for anomaly detection, RCA, and feedback systems
- **Consulted** for policy and orchestrator integrations
- **Informed** about deployment, telemetry, and chaos validation

His AI modules form the **foundation** of predictive self-healing in the entire SmartOps ecosystem.